


The most famous AI coding benchmark stopped being reported this year.


The models had seen the answers.



63%


Cursor audited a top coding agent. In 63 percent of its successful runs it wasn't solving the problem. It found the known fix.

You already own a benchmark nobody trained on. Your git history.



Every bug fix where someone also changed a test is an exam question. The answer never existed on the public internet.

retrospec, run on a real repo.



```
$ python3 retrospec.py mine ~/src/Auto-GPT --since "2023-01-01" \
  --test-cmd "python -m pytest"

Mined 8 case(s) into retrospec-out

$ ls retrospec-out
cases.jsonl  suite.json  patches/

$ jq -r .id retrospec-out/cases.jsonl
92009ceb32ec-more-graceful-browsing-error-handling-3494
aa3e37ac149c-fix-memory-by-adding-it-only-when-context-window-full
20ef13034198-add-tests-for-code-shell-execution-improve-config
5ce6da95fcd0-make-y-n-configurable-3178
adbb47fb652a-scrape-text-regression-test-3387
7cd76b8d8ecb-add-makedirs-to-file-operations-3289

# every case: a real fix from YOUR history.
```

Walks the history, finds the commits that touched source and tests, writes them out as eval cases for agent-eval-harness.

HONEST LIMITS

It's a starting point, not a lab.

It needs a test suite that still runs at older commits.

A commit message is a worse task description than a good ticket.

Runs local. Your code never leaves your machine.

If you had to pick an AI coding model next week, how would you test it on your code?

Not on someone's leaderboard.