

The most famous AI coding benchmark stopped being reported this year.

The models had seen the answers.

63%

Cursor audited a top coding agent. In 63 percent of its successful runs it wasn't solving the problem. It found the known fix.

You already own a benchmark nobody trained on.

Your git history.

Every bug fix where someone also changed a test is an exam question.

Commit message. The task.

Code right before the fix. The starting point.

The tests. The grader.

And the answer never existed on the public internet.

retrospec

Walks the last year of history, finds the commits that touched both source and tests, writes them out as eval cases.

Plugs straight into agent-eval-harness. Run any coding model against your own code and get a number you own.

```
$ python3 retrospec.py mine ~/src/my-project \  
    --since 2025-01-01 --test-cmd "python -m  
pytest"  
Mined 8 case(s) into retrospec-out  
$ ls retrospec-out  
cases.jsonl  suite.json  patches/
```

HONEST LIMITS

It's a starting point, not a lab.

It needs a test suite that still runs at older commits.

A commit message is a worse task description than a good ticket.

Runs local. Your code never leaves your machine.

If you had to pick an AI coding model next week, how would you test it on your code?

Not on someone's leaderboard.